



# Advancing Diabetes Management with Conditional Generative Modeling

Matthew Manion & Sathvika Ayyappa Prabhu  
Chemical Engineering & Computer Science Engineering  
CSE 598: AI For Science  
Fall 2025

# Diabetes: Process Control for Life or Death



Image: CDC

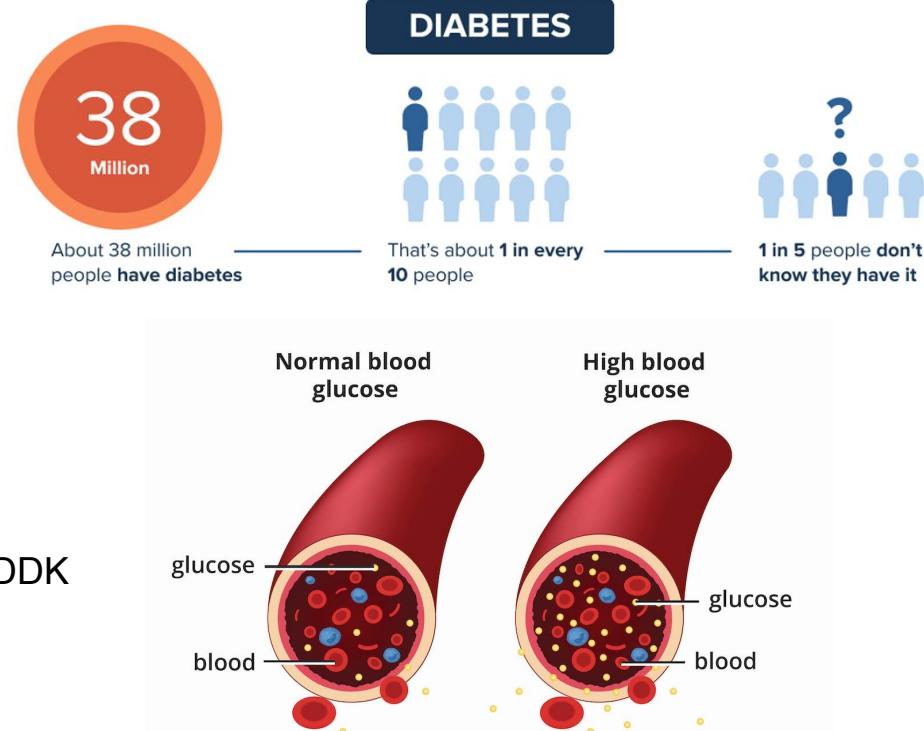


Image: NIDDK

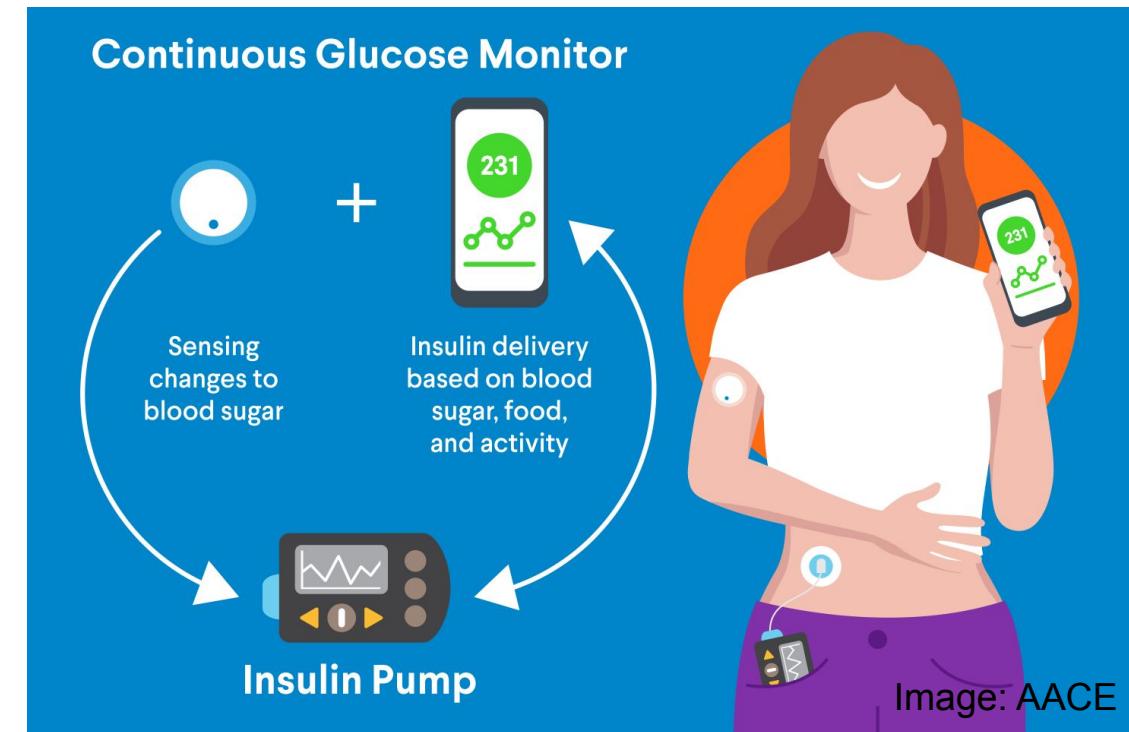


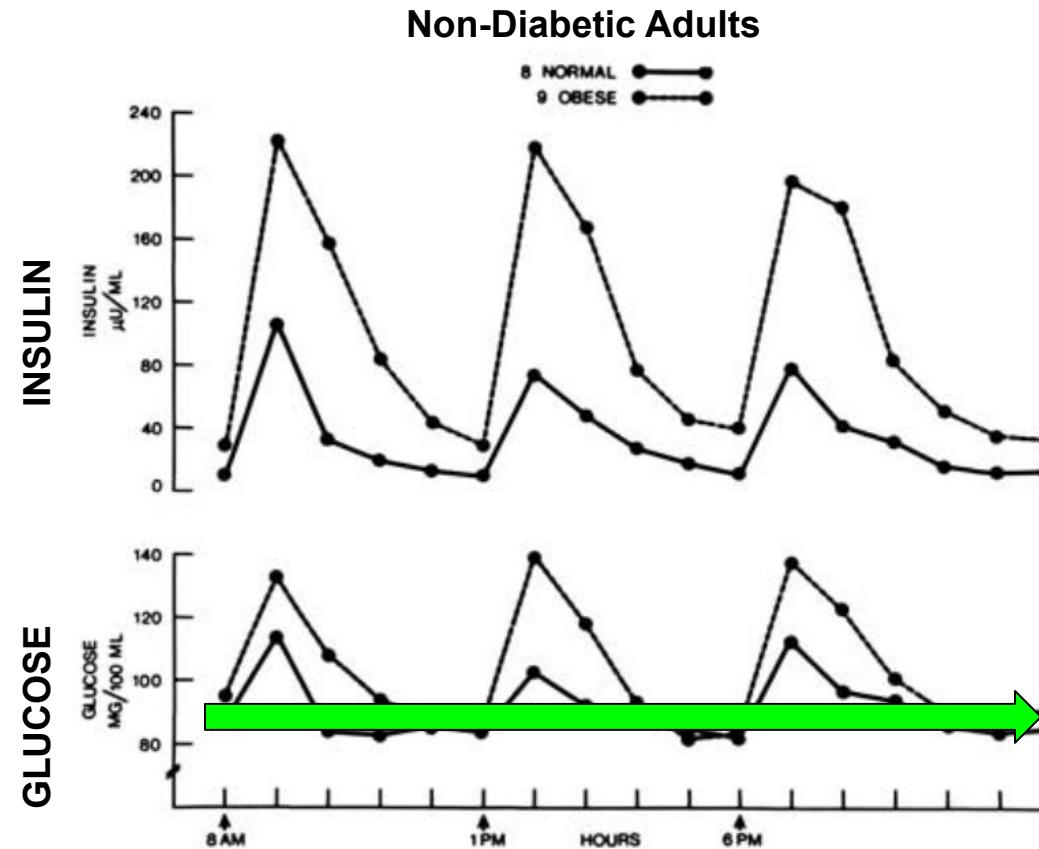
Image: AACE

## Risks of Pump/Monitor Malfunction:

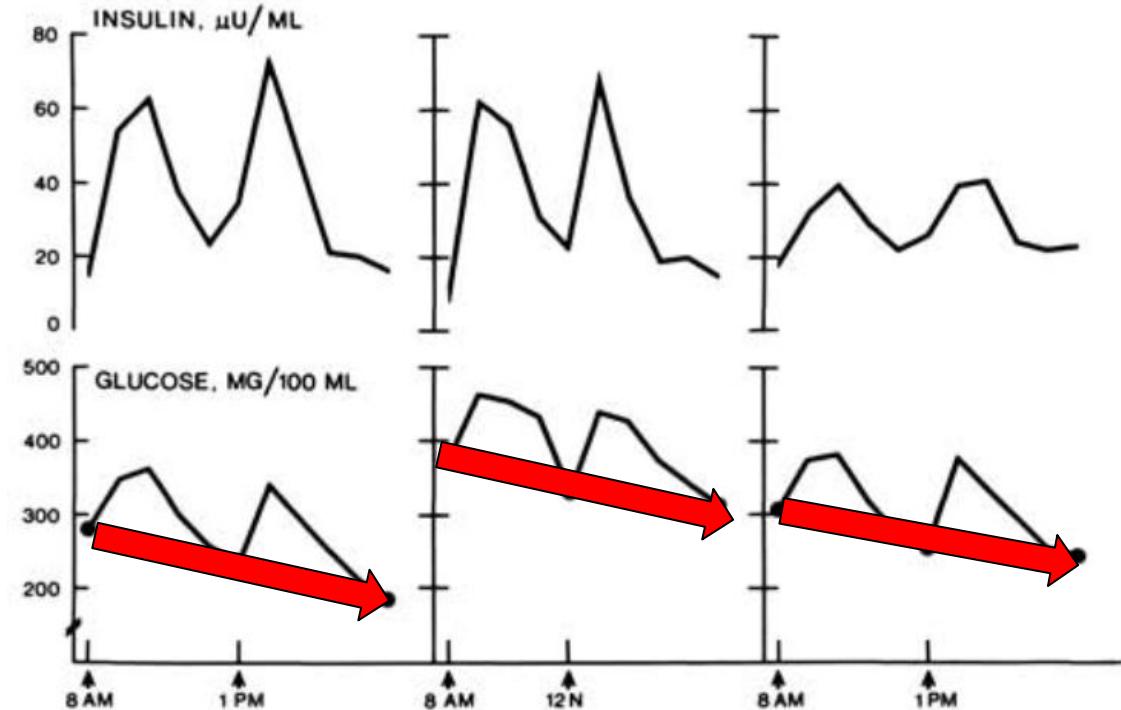
**High Blood Glucose: Coma, Organ Failure**

**Low Blood Glucose: Coma, Seizure, Death**

# Problem: Calculate Insulin Dose for Glucose Trajectory



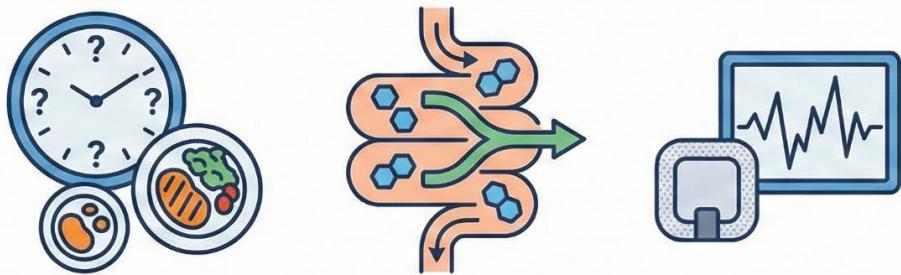
3 Diabetic Adults Given Normal Insulin Dose



SAUL M. GENUTH. Plasma Insulin and Glucose Profiles in Normal, Obese, and Diabetic Persons. Ann Intern Med. 1973;79:812-822

# Blood Glucose Dynamics Are Inherently Stochastic

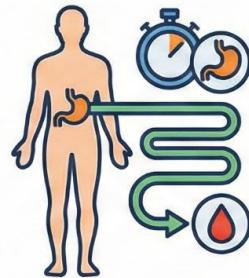
Stochasticity comes from:



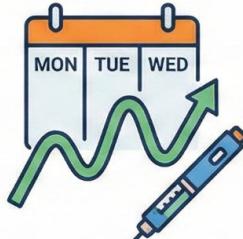
Random Meal  
Timing & Size

Variable Carb  
Absorption

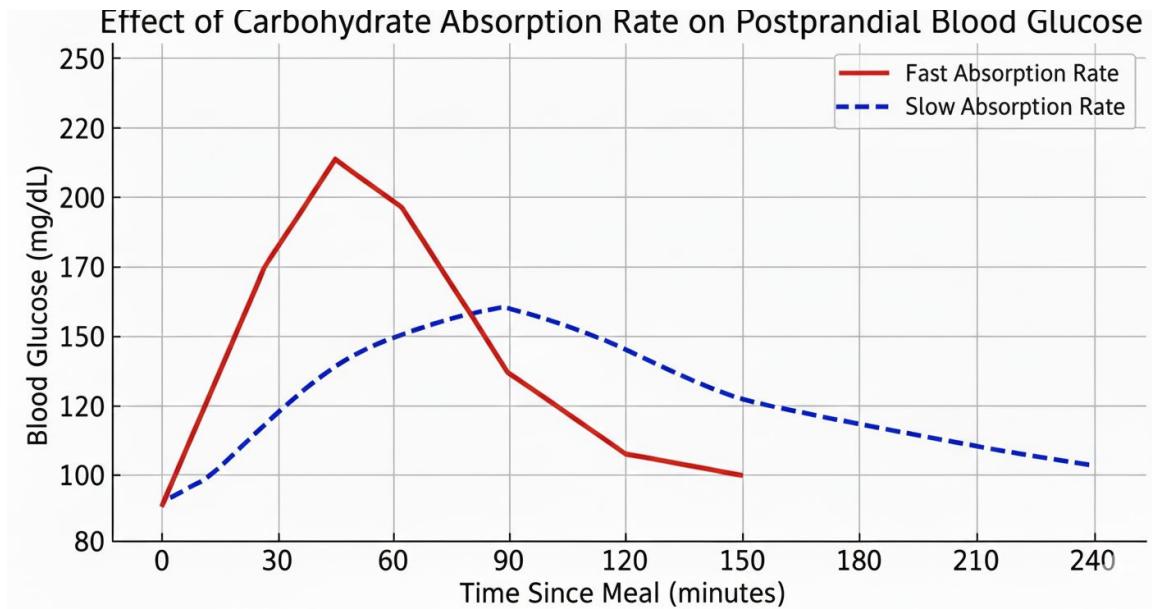
Sensor Noise



Large Physiological  
Delays



Daily Insulin Sensitivity  
Fluctuations



Same observed state → **multiple safe insulin doses**  
Same action → **different BG outcomes**

# Diffusion Policies in Stochastic Offline RL

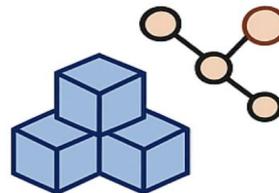


- Expressive
  - Data-anchored
- very successful in deterministic control tasks

However, it is unclear whether they remain:



- Safe
  - Stable
  - Robust
- when transitions are highly stochastic and actions are multimodal



Python implementation of the FDA-approved UVA/Padova Type-1 Diabetes Simulator (2008)



Includes 30 virtual patients  
10 adults, 10 adolescents, 10 children



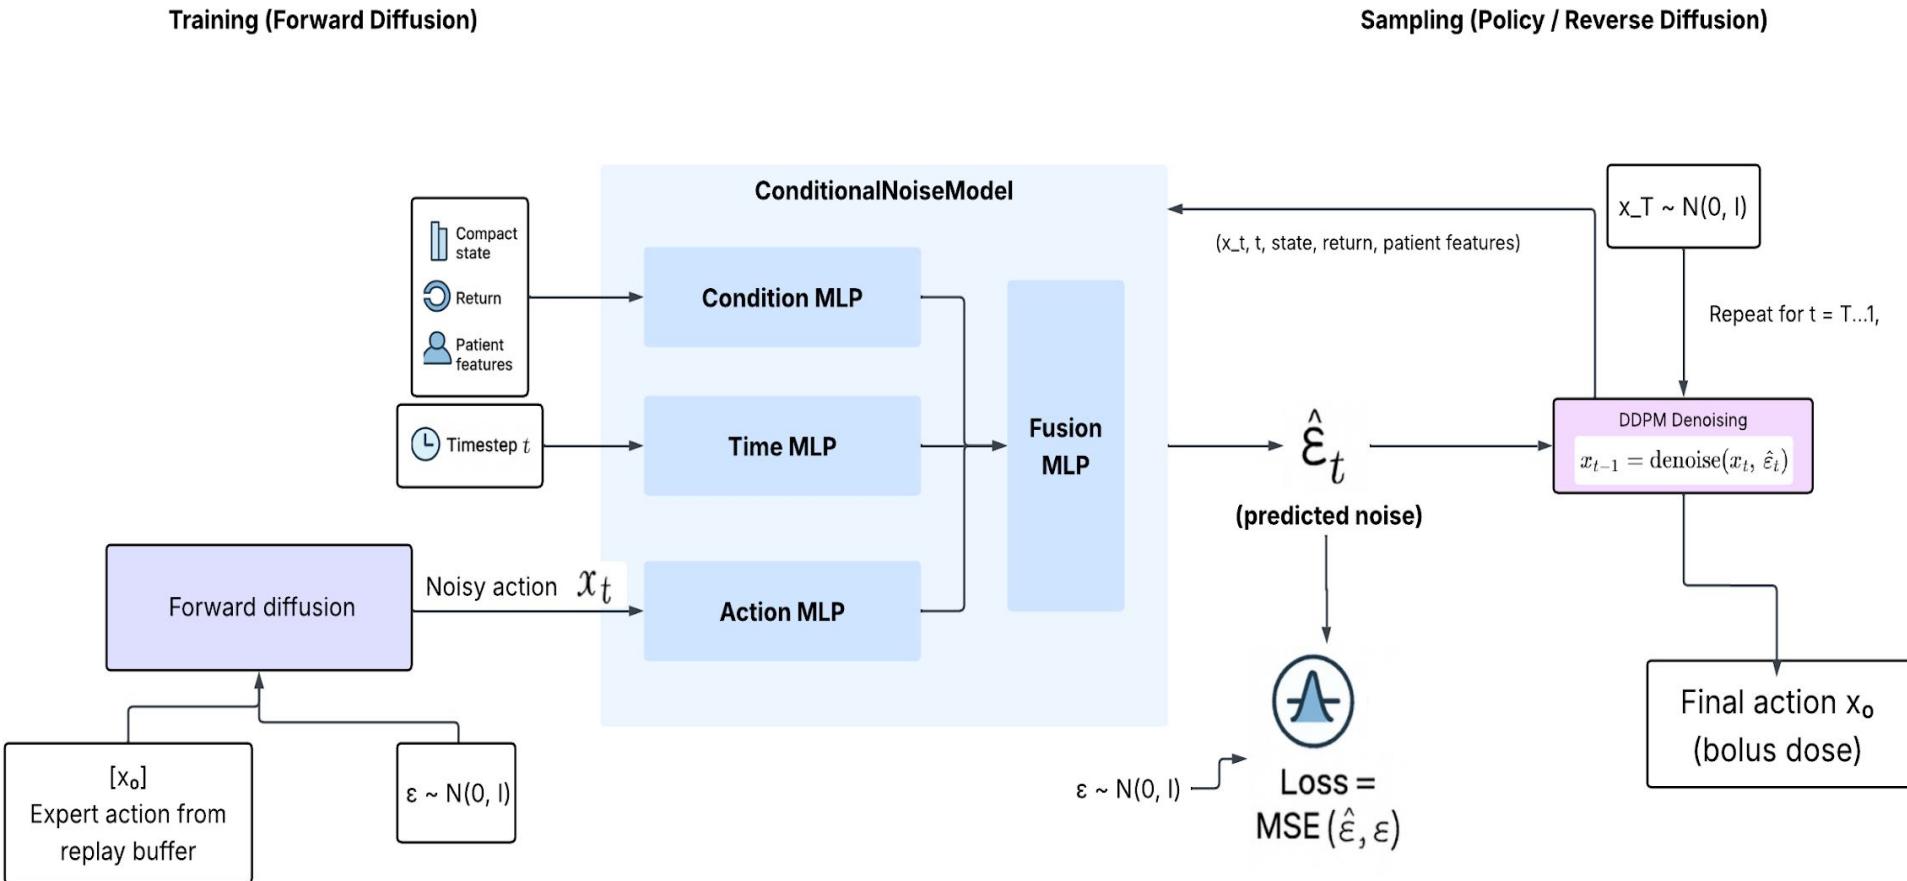
Designed to be “reinforcement-learning-ready”:  
OpenAI Gym / Gymnasium API  
Each step returns (observation, reward, done, info)

We ask:



Can conditional diffusion models reliably handle stochasticity in offline RL?

# Model Architecture



# Problem Definition: Offline RL for Stochastic Blood Glucose Control

We model insulin dosing as a Markov Process Described by:

$$\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{P}, r, \gamma),$$

- **State  $s_t$ :** patient glucose/IOB/CHO context
- **Action  $a_t$ :** Continuous insulin bolus (administered dose)
- **Stochastic Transitions** (meals, absorption, noise, delays) :  $\mathcal{P}(s' | s, a)$
- **Reward  $r_t = -\text{RiskIndex}(G_t)$**  (penalizes unsafe glucose)
- **Policy :**  $\pi_\theta(a|s)$

# Reward Function — Clinical Risk Index

We use the clinical Blood Glucose Risk Index, which quantifies how dangerous a glucose value is.

Risk is split into two components:



**LBGI (Low BG Risk Index)**



**HBGI (High BG Risk Index)**

Hypoglycemia carries much higher risk, so LBGI is weighted more heavily.

For each glucose value  $G_t$ , we compute a risk score:

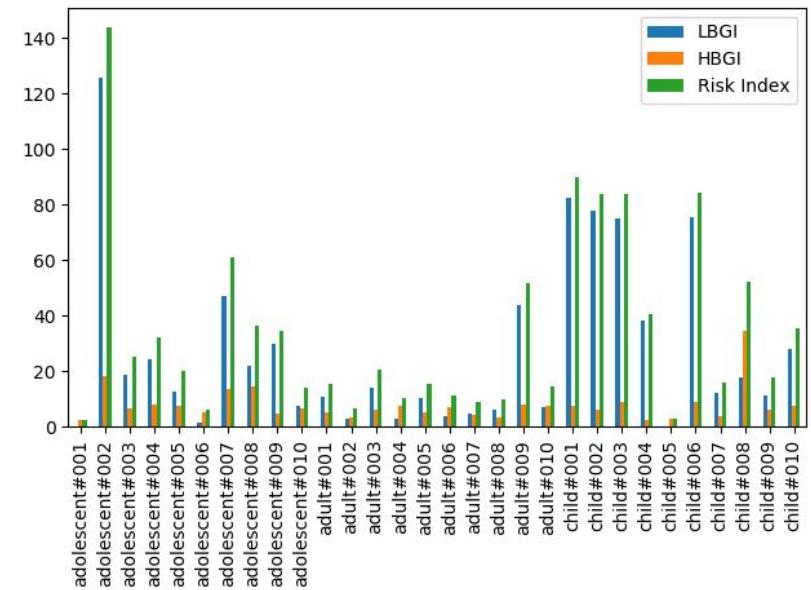
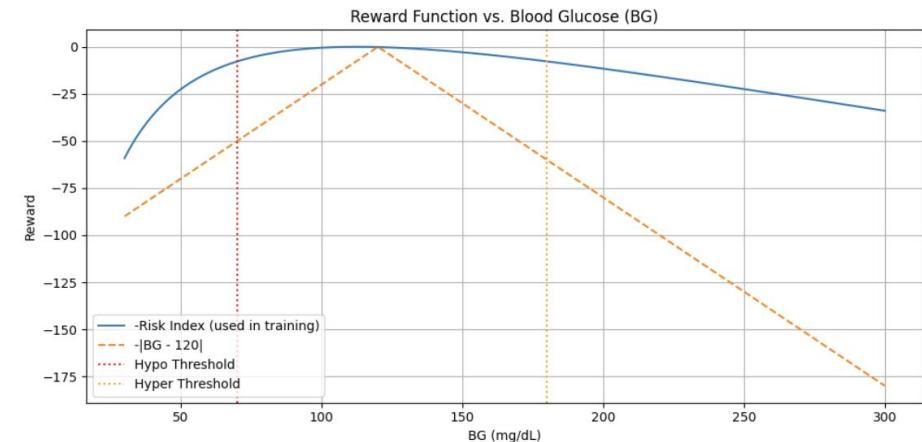
$$\text{Risk}_t = \lambda \cdot \text{LBGI}_t + (1-\lambda) \cdot \text{HBGI}_t, \lambda > 0.5$$

Reward is:



=

$$r_t = -\text{Risk}_t$$



# Diffusion Model as a Policy $\pi(a_t | s_t)$



We treat an insulin bolus  $a_t$  as a sample generated from noise



Forward process:  
add Gaussian noise to dataset actions  
→ produce noisy samples

Reverse process: learn to denoise step-by-step conditioned on:

- current state  $s$
- patient physiology
- target return
- diffusion timestep  $t$

At test time:

Start from pure noise

Iteratively denoise

→ final action = insulin bolus



## Forward (training):

Noise injection to create supervision:

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

- $x_0 = a_t$  (true bolus from offline dataset)
- Produces noisy action  $x_{txt}$  at diffusion step  $t$ .

## Reverse (policy):

Model predicts noise:

$$\hat{\epsilon}_\theta(x_t, s_t, t, \text{patient}, \text{target})$$

Used to denoise:

$$x_{t-1} = \text{Denoise}(x_t, \hat{\epsilon}_\theta)$$

## Training objective (Noise Prediction Loss):

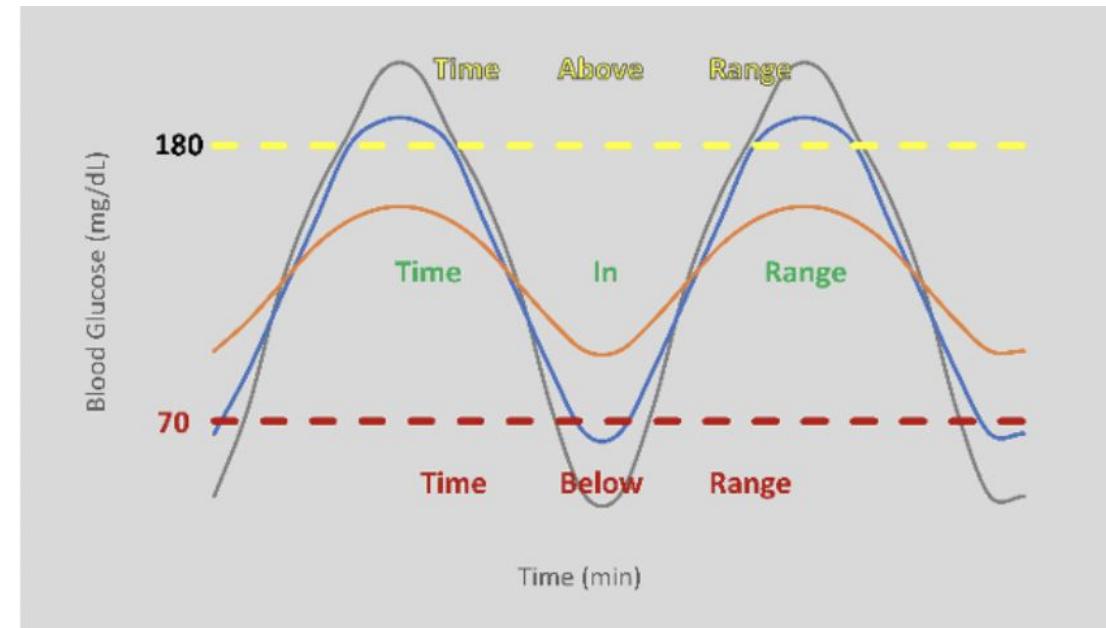
$$\mathcal{L}(\theta) = \|\epsilon - \hat{\epsilon}_\theta(x_t, s_t, t, \text{context})\|^2$$

# Experimental Setup & Evaluation Protocol

- **Environment:** SimGlucose ([T1DSimEnv](#), adult#001, RandomScenario)
- **Resolution:** 5-minute steps
- **Episode length:** 24 hours → 288 control decisions per episode
- **Training run:** 600 episodes (~600 virtual days)
  - 75 expert-only pretraining episodes (BBController)
  - 525 episodes with diffusion + expert mixing
- **Oracle-style aggregation:**
  - Start with pure Expert trajectories
  - Then mix in BB vs diffusion actions
    - i. Model attempts to predict noise
    - ii. If TIR <55%, Model decisions not saved
    - iii. Expert is given same situation, decisions saved for training
  - All transitions stored in a replay buffer (1M capacity)
- **Optimization:**
  - Batch size = 256
  - 5 diffusion updates per episode
  - Cosine LR schedule (3e-4 → 1e-6 across 600 episodes)
- **Policies evaluated:**
  - **BBController** (Behavior policy / data source)
  - **PID controller**
  - **Tabular Q-learning**
  - **TD3-BC**
  - **Diffusion Policy (ours)**
- **Offline constraint:**
  - No environment interaction during training
  - All models trained only on the Expert & Model replay buffer

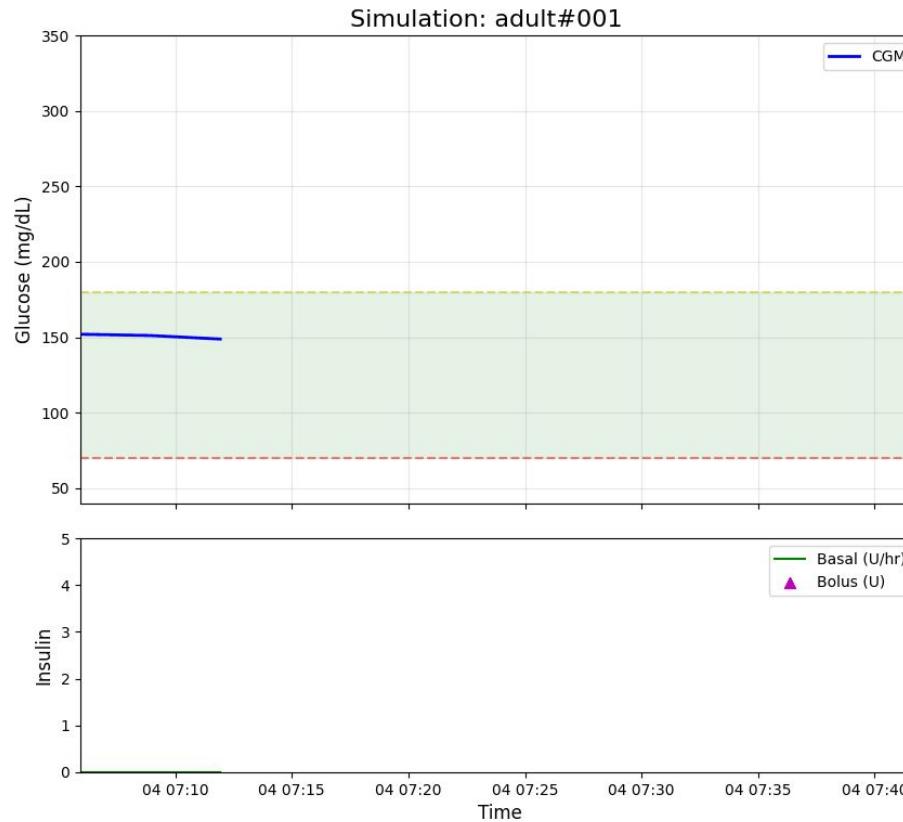
# Clinical Evaluation Metrics

Metric	Definition
TIR (%) Time in Range	% of time BG is <b>70–180 mg/dL</b>
TBR (%) Time Below Range	% of time BG is <b>&lt; 70 mg/dL</b> (hypoglycemia)
TAR (%) Time Above Range	% of time BG is <b>&gt; 180 mg/dL</b> (hyperglycemia)
CV (%) Coefficient of Variation	Glucose variability ( <b>SD ÷ mean × 100</b> )

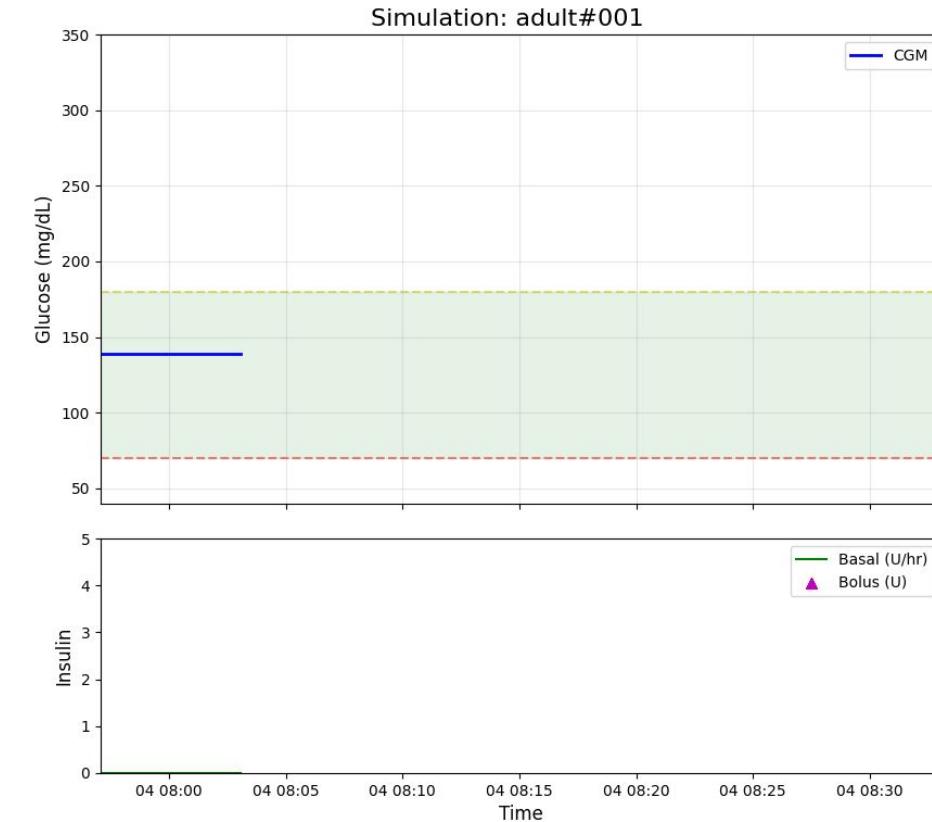


# Example Diffusion Training

Early Training (Patient Dies)



Late Training



# Results: Diffusion Trained Only on Body Weight

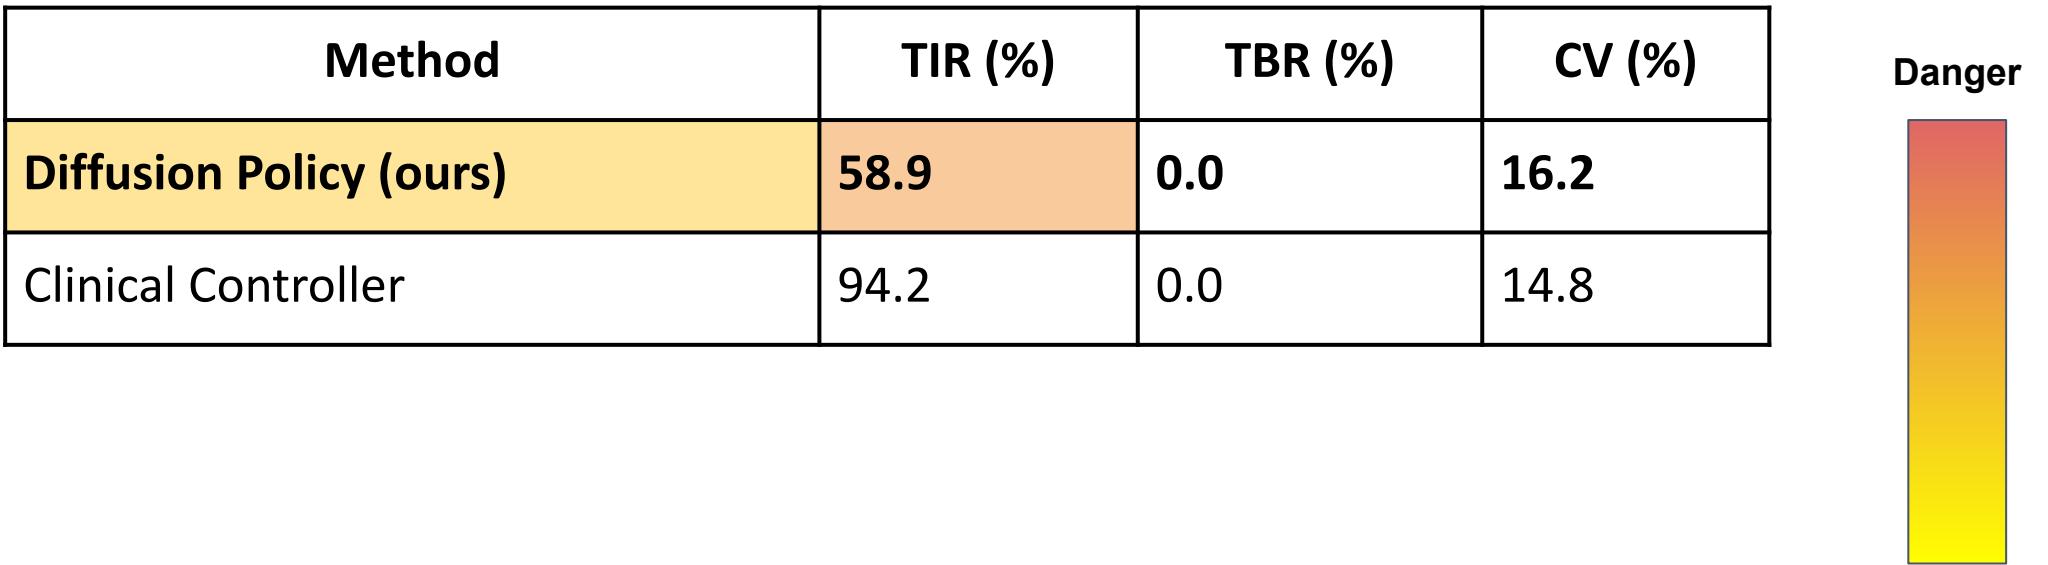
Method	TIR (%)	TBR (%)	CV (%)	Danger
PID Controller	68.2	6.1	32.1	
Tabular Q-Learning	70.1	5.2	30.3	
TD3-BC	71.8	4.5	28.4	
<b>Diffusion Policy (ours)*</b>	<b>100.0</b>	<b>0.0</b>	<b>13.35</b>	
Clinical Controller*	91.3	0.0	8.83	

\*Averaged Over 3 Adult Evaluations

# Results: BW Generalization to Other Groups

Method	TIR (%)	TBR (%)	Danger
<b>Diffusion Policy (ours) (Adolescent 1)</b>	79.5	8.3	
Clinical Controller (Adolescent 1)	71.2	0.0	
<b>Diffusion Policy (ours) (Child 1)</b>	49.0	51.0	
Clinical Controller (Child 1)	76.0	12.8	

# Results: Diffusion Trained on Multi Params



Averaged Over 3 Adult Evaluations

# Key Takeaways & Next Steps:

- Diffusion policies are effective for stochastic offline RL in blood glucose control.
- Performs reliably on adult virtual patients, where physiological dynamics are more stable.
- Generalization to adolescents/children remains challenging due to higher variability and noisier glucose–insulin dynamics.
- Training diffusion on multiple patient parameters requires significant compute, especially under stochastic transitions.

Next:

- Scale up compute
- Change Reward Scheme: Emphasize Ideal Glucose Region
- Expand Time Horizon

# Questions?

# The RL Opportunity

- RL can learn personalized insulin policies.
- It can adapt to:
  - Each patient's physiology
  - Time-of-day patterns
  - Meal-driven disturbances
- But online RL is risky in healthcare : unsafe, unethical, expensive
- → **Offline RL** is the only viable option.

# Data Collection



## Environment: T1DSimEnv

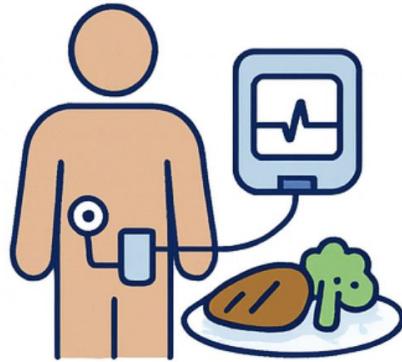
(simglucose, FDA-approved UVA/Padova simulator)



## Behavior policy:

Built-in Basal-Bolus  
(BBController)

Represents standard clinical dosing behavior



## Stochastic meal scenarios



Randomized meal timing, meal size, and absorption variability



## Physiological variability

Adult virtual patients have diverse glucose/insulin kinetics

- **Dataset generation:**
  - Run many full-day simulations per adult virtual patient
  - Log transitions:

$(s_t, a_t, r_t, s_{t+1}, \text{info})$
  - ( $\approx 300\text{--}400$  episodes  $\rightarrow \sim 7.5 \times 10^5$  transitions)

# Forward & Reverse Processes

## Forward (training):

Noise injection to create supervision:

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

- $x_0 = a_t$  (true bolus from offline dataset)
- Produces noisy action  $x_{txt}$  at diffusion step  $t$

## Reverse (policy):

Model predicts noise:

$$\hat{\epsilon}_\theta(x_t, s_t, t, \text{patient}, \text{target})$$

Used to denoise:

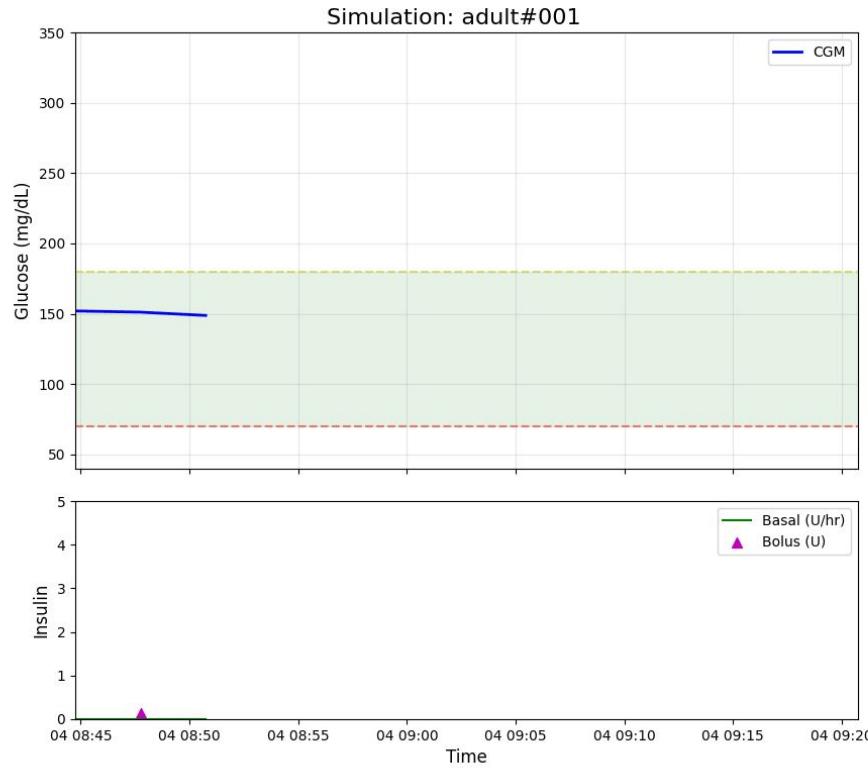
$$x_{t-1} = \text{Denoise}(x_t, \hat{\epsilon}_\theta)$$

## Training objective (Noise Prediction Loss):

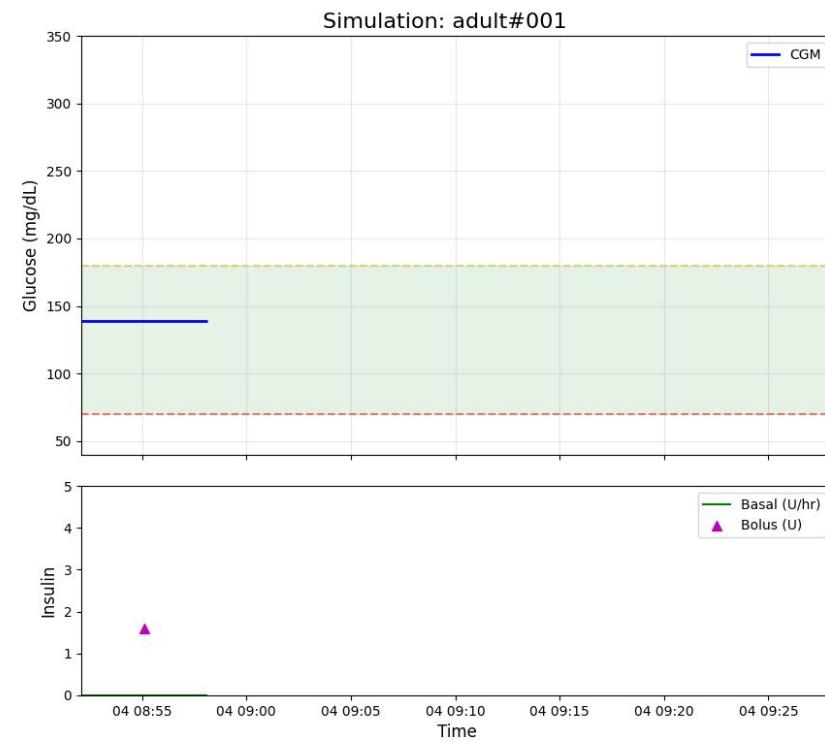
$$\mathcal{L}(\theta) = \|\epsilon - \hat{\epsilon}_\theta(x_t, s_t, t, \text{context})\|^2$$

# Example Diffusion Training: Latest

Early Training (Patient Dies)



Late Training

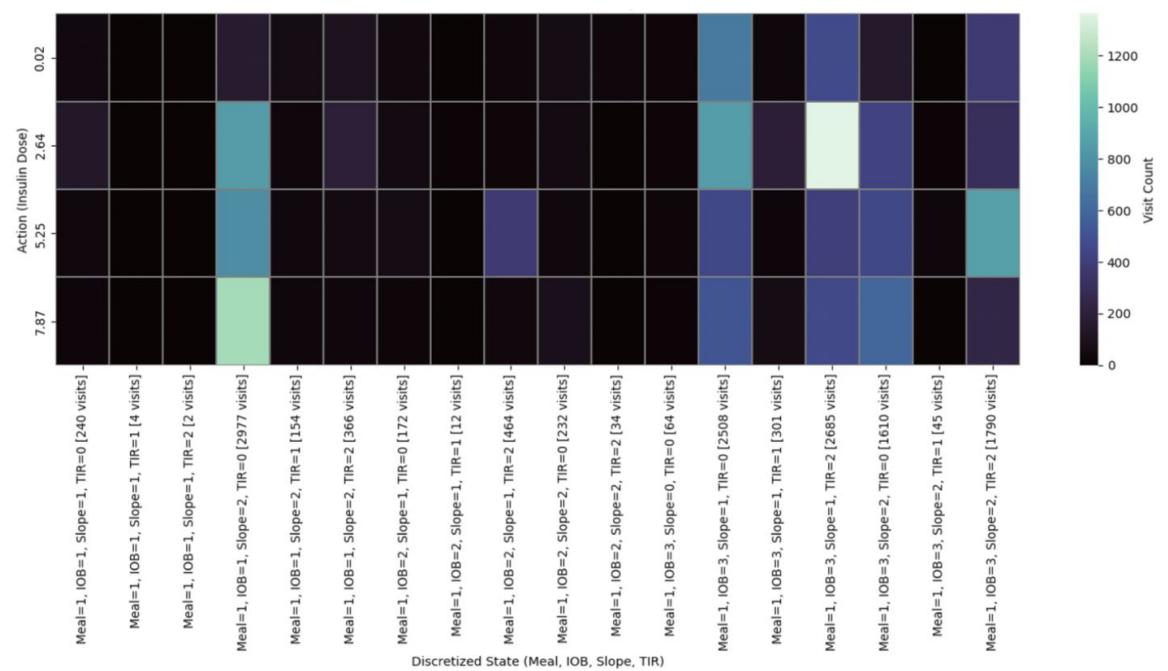
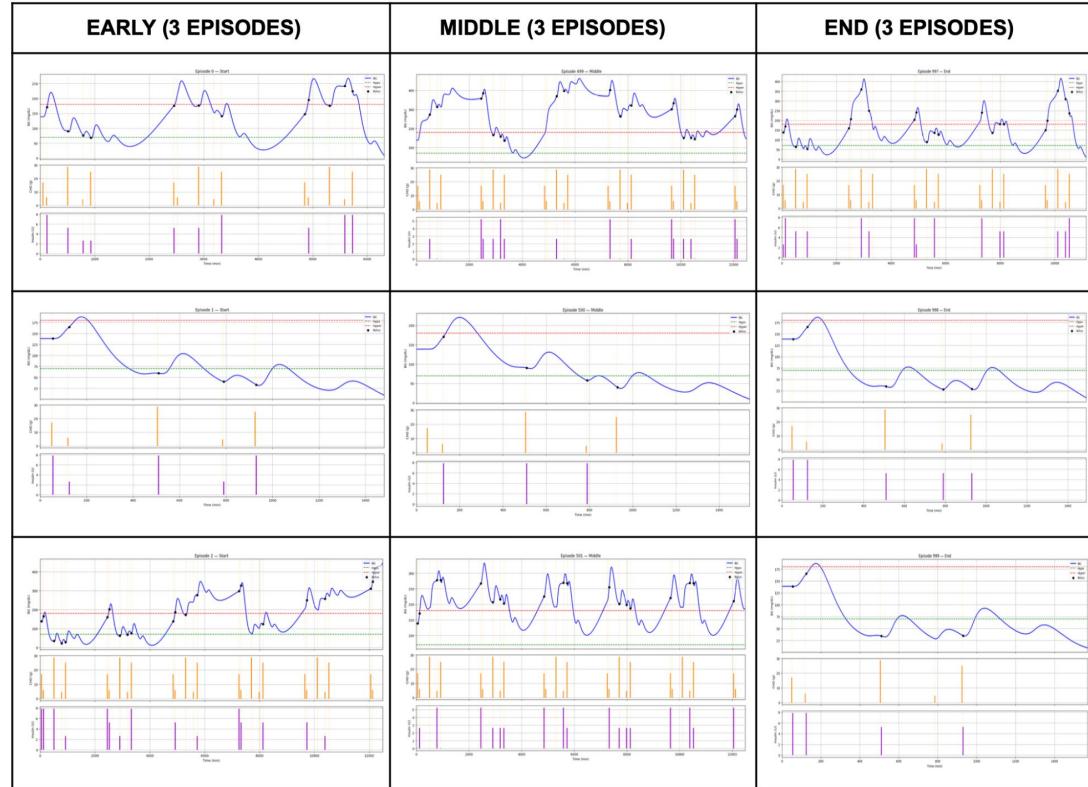


# Model Architecture

Placeholder..

Diagram comes here. Will remove/update text  
after.

# Results: Comparison Across Controllers (TAB RL)



# Stochasticity Creates a General Offline RL Challenge

- Offline RL must learn from a fixed dataset → no corrections
- Under stochasticity:
  - a. Q-values become unreliable
  - b. Policies collapse multimodal distributions into unsafe averages
  - c. Models produce out-of-distribution actions

**How do we design policies that remain robust under stochastic offline RL?**